

CliffGuard: Degeneration-Aware Scoring Reverses the Apparent Refusal Trend Under RTN Quantization

Trilochan Sharma

Independent Researcher

trilochan2625@student.ku.edu.np

August 6, 2026

Abstract

Post-training quantization is widely reported to erode refusal behaviour. We argue that a large part of that report is an artifact of how refusal is scored, and we introduce **CliffGuard**, a degeneration-aware paired protocol that separates three outcomes a conventional phrase-matching pipeline collapses into two: refusal, compliance, and degenerate output. Applying it to a round-to-nearest ladder from 8.5 down to 2.5 stored bits per parameter, with 500 paired prompts per rung and a 7B judge behind the degeneracy gate, reverses the expected direction of the effect. Within the band where output stays coherent, quantization makes both models tested *more* likely to decline — where declining is the judge’s class of refusing, deflecting, redirecting or warning, which we characterise rather than assume — at 1.15 points of that class per bit removed (95% CI [0.75, 1.55], prompt-level cluster bootstrap). Under the 7B judge this reaches significance at 4.5 bits on both models ($p_{14} = 0.021$ and 0.011), but re-grading the identical completions with three graders from two other families does *not* reproduce it: three of the four resulting model $\times$ grader comparisons agree on the direction, none reaches significance, and pairwise agreement with the 7B judge is 64.9–78.4%. We withdraw the significance claim and report only the direction. A dependent variable that moves by a factor of thirty with the choice of grader is this paper’s own thesis applied to itself. Transitions from full-precision refusal to coherent compliance never exceed 2.2%, and a simultaneous one-sided 95% bound puts every rung at most 4.62%. Scoring the identical completions with a refusal-phrase list instead reports up to 38.4%, and still differs by a factor of 4.9 at a rung where no completion is flagged as degenerate. The phrase-list estimator’s two paired factors move in opposite directions as the marker list grows, which removes any structural guarantee that enlarging the list moves the estimate toward a target — and on these runs it is non-monotone in fact. We separate the remaining gap into a gate effect and a grader effect that dominate at different rungs. Two failures are invisible to safety-oriented instrumentation: a frozen refusal probe, fitted on the judge’s own labels so both arms share a target, retains 95–100% of its full-precision discriminability across 8.5–4.5 bits, spanning both models’ 4.5-bit refusal shifts, and arithmetic accuracy degrades at a threshold differing by a full bit between model families. Every refusal result is a label assigned by a model judge to a greedy 48-token completion against each model’s own full-precision baseline: a behavioural *change* measure, not an absolute harmfulness rate. We close by stating, as a testable design hypothesis, what these measurements imply for how a quantizer should be selected.

1 Introduction

Post-training quantization [Frantar et al., 2023, Lin et al., 2024, Dettmers et al., 2023] is how large language models reach consumer hardware, and a body of work reports that it damages safety alignment [Hong et al., 2024, Egashira et al., 2024, 2025, Kadadekar, 2026], motivating alignment-preserving quantization methods [Wee et al., 2025]. The practical reading is that a model validated at full precision may answer harmful requests once compressed.

We set out to characterise that effect. On the models and prompts studied here it does not appear when the refusal decision is read off generated text. What appears instead is a shift in the opposite direction, consistent in sign under every grader we tried, together with a capability loss that no safety-oriented instrument in our battery detected. Whether that shift reaches significance turns out to depend on which grader is asked; §4 reports the test that established this and withdraws the claim it defeats.

Findings.

1. Inside the band where output stays coherent, lower precision produces *more* refusal, at 1.15 points per bit removed, 95% CI [0.75, 1.55] (§4.1). The relationship is not linear across the whole ladder: from FP16 down to 8.5 bits refusal moves by +0.4 and +0.2 points. Both the slope and its interval come from the 7B judge’s labels, so the interval prices sampling error and not the grader dependence that finding 6 documents.
2. Refusal-to-compliance transitions stay at or below 2.2% at every rung on both models; a simultaneous 95% upper bound over all 14 model $\times$ rung cells is 4.62% (§4).
3. Phrase-list scoring of the same completions reports up to 38.4%. The error separates into a gate effect and a grader effect that dominate at different rungs, and the estimate is *non-monotone* in the phrase list, so enlarging that list carries no structural guarantee of improving it (§5).
4. Arithmetic capability collapses at a model-specific bit-width, one full bit apart across families (§6).
5. A frozen refusal probe holds 95–100% over 8.5–4.5 bits, spanning both models’ 4.5-bit refusal shifts. It falls to 57% at Phi-3.5-mini’s largest shift, at 3.5 bits. Its full-precision d' is only 0.74 on Qwen2.5-3B, so what it retains there is weak to begin with (§7).
6. Re-grading the identical completions with three graders from two other families reproduces the *direction* on three of the four model $\times$ grader comparisons and the significance on none, at 64.9–78.4% pairwise agreement with the 7B judge (§4). We withdraw the significance claim. This is the paper’s own thesis applied to itself, and it is why the word “significant” appears nowhere above.

What this paper does not measure. The prompt set carries no per-prompt harmfulness annotation, so a transition from baseline refusal to quantized compliance is exactly that — a change in the model’s own decision — and not an observation of harmful content. We use the term *refusal-to-compliance transition* throughout and avoid “harmful compliance”. §9 states what would be needed to license the stronger term.

1.1 The CliffGuard protocol

The results above are not separable from how they were measured, so we state the measurement procedure as a protocol rather than as a list of implementation details. CliffGuard is that protocol: an evaluation method, not a quantizer, a model, or a released benchmark. Its defining commitment is that *degeneration is a third outcome, decided before refusal*, because the conventional binary rule

$$\text{contains a refusal phrase} \longrightarrow \{\text{refusal, compliance}\}$$

assigns incoherent output to *compliance* by default, and on each RTN ladder we tested the lowest rung is exactly that.

1. **Pair.** Generate completions for the same prompts from the full-precision checkpoint and from each candidate quantization, holding decoding fixed. All comparisons are within-prompt.

2. **Gate.** Decide degeneration first and independently of refusal, using more than perplexity (§3.1). Report the degeneration rate beside every other rate.
3. **Grade.** Label surviving completions semantically into {refusal, compliance, unclear}, keeping UNCLEAR as its own state rather than folding it into either decision.
4. **Transition.** Report the two directions separately, refusal to compliance and compliance to refusal, with the number of gradable pairs, paired tests, and multiplicity correction. A single marginal refusal rate hides which prompts moved and which way.
5. **Stress the scorer.** Re-score identical completions under alternative marker lists and under an alternative gate, and publish the spread. The marker list is an experimental variable, not an implementation detail.
6. **Separate capability from fluency.** Score a gold-labelled task with no classifier in the loop, so capability loss can be observed independently of both the grader and the degeneration detector.

Steps 2 and 4 are what the rest of this paper turns on: without step 2 the phrase-list estimate at 2.5 bits is 38.4% instead of 0.2% (§5), and without step 4 the conservative shift that dominates our results is invisible (§4). Step 5 is what exposes the structural defect of Equation 3. The protocol’s own weak point is that it presumes a trustworthy semantic grader, and validating ours against blinded human labels is the largest piece of work it still needs (§9).

2 Related work

Quantization and safety. Hong et al. [2024] evaluate trustworthiness of compressed models across dimensions including safety. Egashira et al. [2024] and Egashira et al. [2025] show that quantization can be *adversarially* exploited: a model benign at full precision can be constructed to become harmful once quantized. Our question is different: whether ordinary, non-adversarial round-to-nearest quantization degrades refusal on its own.

Kadadekar [2026] is the closest recent point of comparison and reaches an apparently opposite conclusion: across six models on a GGUF ladder plus AWQ and GPTQ checkpoints, that work reports rows where quality is stable or improved while refusal falls by 12–68 points, scored through a refusal-template stability index. We do not claim that result is wrong, and our design cannot adjudicate it: we vary one quantizer family where they vary three, and our prompts, models and grader all differ. We do note that a template-derived refusal score is the estimator whose failure mode we characterise in §5, and that our own phrase-list column reproduces a fall of that order (Table 3) on completions the judge and the degeneracy gate both say are not compliance. Distinguishing a genuine refusal loss from this artifact requires reporting a degeneracy rate beside every refusal rate, which we recommend in §8.

Refusal representations. Ardit et al. [2024] establish that refusal is mediated by a single activation direction, using causal ablation. Zhao et al. [2025] find harmfulness and refusal to be separately encoded. Chhabra and Khalili [2025] study refusal in compressed models mechanistically and report that the refusal direction is largely preserved under compression; our §7 is consistent with that at moderate precision and extends it in one respect, by pairing the preserved direction against a behavioural measurement on the same rungs and showing the two dissociate. Duan [2026] benchmark activation-monitor staleness across model updates and likewise find quantization causes little monitor degradation. Read together with our results, the agreement is the problem rather than the reassurance: a monitor that does not move while behaviour does is not a safety signal. Linear probes are standard [Alain and Bengio, 2017], and Belinkov [2022] catalogues exactly this interpretive limit.

Grading refusal. Automatic graders have moved from string matching toward classifiers and judge models [Inan et al., 2023, Ghosh et al., 2024, Zheng et al., 2023], and benchmark suites ship their own [Mazeika et al., 2024]. Souly et al. [2024] make the argument closest to ours in the jailbreak setting: success rates are inflated when a grader counts non-substantive output as compliance. We find the same failure mode under quantization, where non-substantive output is not an occasional artifact but the terminal state of the ladder.

Degeneration. Holtzman et al. [2020] characterise repetition loops under likelihood-maximising decoding. We use that observation in reverse: because such text is high-likelihood, perplexity cannot detect it, which is what makes a naive degeneracy filter fail (§3.1).

3 Method

Ladder. Per-group asymmetric round-to-nearest, group size 64, applied to every transformer-block linear layer; embeddings and the output head remain at FP16, as practical schemes do. Bit-widths are 8, 7, 6, 5, 4, 3, 2. We label rungs by $b + 32/64$ bits per quantized parameter, the stored cost once the FP16 scale and zero point are amortised. This is a label for the quantized layers only: because embeddings and the head stay at FP16, whole-model bits per parameter are higher, and by an amount that differs across architectures. No claim here depends on the label being a model-wide average; it is an ordinal dose axis.

We use round-to-nearest rather than a mixed-type format because it varies *only* bit-width. Mixed-type formats vary block structure and per-tensor type assignment simultaneously, which would leave any observed threshold uninterpretable. The cost of that choice is external validity, and we state it plainly in §9: nothing here establishes how AWQ, GPTQ or GGUF k-quants behave.

Models and prompts. Qwen2.5-1.5B-Instruct and Qwen2.5-3B-Instruct [Qwen Team, 2024], and Phi-3.5-mini-instruct [Abdin et al., 2024]. Prompts are 500 first-turn human utterances from Anthropic HH-RLHF [Bai et al., 2022], taken as the first 500 deduplicated rows in file order rather than by random draw. This is a fixed benchmark, not a probability sample from a defined population, and it governs how every interval in this paper should be read.

On these 500 prompts the transition rate is not an unknown parameter — it is counted. The bootstrap and the exact binomial bounds are therefore not benchmark-level confidence statements; they answer a different and explicitly assumption-bearing question: if these prompts are treated as exchangeable draws from the first-turn HH-RLHF pool, how much would the estimate move on a comparable draw? That assumption is not testable here, and every “95%” in this paper should be read as conditional on it. Two further conditions apply where noted: the drift intervals are additionally conditional on a coherent band selected from the observed degeneracy rates and then held fixed (§4.1), and the probe ranges are dispersion over splits of one fixed prompt set rather than intervals at all (§7). Completions are greedy, 48 new tokens, left-padded batching. The residual stream, where used, is read at mid-depth resolved from each model’s own configuration.

Labels come from the model. Every rung is compared against the *same model at FP16* on the *same prompts*. We therefore measure behavioural *change* under quantization, and make no claim about absolute safety.

Which models appear in which arm. Qwen2.5-1.5B contributes to the capability and probe results, not to the refusal results. Its completions were graded by a 1.5B judge, which returned REFUSE for 100% of full-precision completions including plainly helpful ones; a grader with no discriminative power cannot support a refusal result. The capability arm scores against gold

arithmetic answers and the probe arm against activations, so neither needs that grader and both include all three models. We flag this because the same saturation is a live risk for any small self-judging setup.

3.1 Degeneracy must be detected before refusal

Completions are labelled **refusal**, **compliance**, or **degenerate**, with degeneracy decided first. At low precision a model emits text such as

```
on \n: .@@\n \ \ \ \4\n ""lllv.: \noo@@\noo@@
```

which contains no refusal phrase. Any classifier that decides refusal versus compliance without a third category scores this as compliance.

We flag degeneracy when any of the following holds: mean per-token negative log-likelihood under the FP16 reference exceeds $3\times$ that reference’s own median; distinct-trigram ratio below 0.60; largest single-token share above 0.35; alphabetic-character fraction below 0.70.

Perplexity alone is insufficient, and it fails in the direction that matters. Repetition loops are high-likelihood by construction [Holtzman et al., 2020]. On Qwen2.5-3B the 2.5-bit rung’s median NLL is 4.13, *lower* than the partially-coherent 3.5-bit rung’s 5.80. A perplexity-only filter ranks the two rungs backwards, admitting the more degraded output and rejecting the less degraded one: at 2.5 bits it flags 15.6% of completions as degenerate where the composite detector flags 99.8%.




Figure 1: Left: median negative log-likelihood is not monotone in bit-width, so a perplexity-only degeneracy filter orders the lowest two rungs incorrectly. Right: distinct-trigram ratio, largest single-token share and alphabetic fraction separate coherent from degenerate output without ambiguity. Qwen2.5-3B.

The thresholds are handcrafted, so we report what they do rather than asserting they are right. We have no human-labelled degeneration set, and without one the four cut-offs cannot be called calibrated. Three measurements bound how much that matters. *Where the result lives, the gate is inert.* At 4.5 bits and above, not one of the four rules fires on a single completion of either model. The headline is therefore untouched by any threshold choice.

The rules are not redundant, and the perplexity rule is the weak one. At Qwen2.5-3B’s 2.5-bit rung the alphabetic-fraction rule fires on 99.6% of completions and the largest-token-share rule on 81.8%, while the NLL rule fires on only 15.6%. Deleting the NLL rule entirely leaves that rung at 99.6% degenerate. This is the repetition-loop failure of §3.1 quantified: the likelihood term is close to useless exactly where degeneration is worst, and the surface statistics carry the decision.

The coherent band does not move. Nudging all four thresholds by 20% simultaneously — in either direction — leaves both models’ coherent bands exactly as reported, {4.5, ..., 8.5} for Qwen2.5-3B and {3.5, ..., 8.5} for Phi-3.5-mini. The band selection that Equation 1 conditions on is not balanced on a threshold.

What none of this establishes is that a completion the gate calls degenerate is one a human would. That needs the labelled set (§9).

The gate is a confound, and we bound it rather than assume it away. Gating on likelihood under the FP16 reference before grading is outcome-dependent: a coherent answer that full precision finds improbable is exactly what a genuine quantization-induced transition would look like, and it is also what the NLL criterion rejects. Three facts bound the damage. First, the likelihood term is *unconditional*: the reference scores each completion on its own, never given the prompt that produced it. It therefore measures whether the text is fluent under the reference model, not whether it is a probable answer to that request, and it cannot select on the prompt–answer relationship that a safety transition consists of. A coherent answer to a harmful request is ordinary English and scores like ordinary English. Second, at every rung above 4 bits the gate removes nothing: degeneracy is 0.0% for both models from 8.5 through 4.5 bits, so the transition rates in Table 1 over that range — which carries the headline results at 4.5 bits and above — are effectively ungated, and the two gates admit an identical set of completions there. Phi-3.5-mini’s largest shift, at 3.5 bits, falls outside it, at a rung where the gate removes 0.6%. Third, where the gate does bind, we vary it rather than fix it: Table 3 reports the same phrase list under both gates, which is how the gate’s contribution is separated from the grader’s in §5. The gate is load-bearing only at 3.5 bits and below, those rows are reported with their effective sample sizes rather than folded into a headline, and a blinded audit of the completions it removes is not yet done.

Grading. Non-degenerate completions are graded by Qwen2.5-7B-Instruct [Qwen Team, 2024] loaded in 4-bit [Dettmers et al., 2023]. The judge is shown the prompt and the completion and asked whether the assistant complied or refused, with an explicit instruction that declining, deflecting, redirecting or warning counts as refusal in any language, and a third UNCLEAR option for incoherent or unrelated text. The decision is read from the first-token logits over the three label words rather than from generated text, which makes it deterministic and removes any dependence on the judge staying on format when shown degraded input. Each label is two tokens under this tokenizer, but their *first* tokens are distinct ($\square_{\text{REF}}$, $\square_{\text{COM}}$, $\square_{\text{UNC}}$; ids 38029, 7682, 75255), so the three-way argmax is well defined; the runner asserts this at load time rather than assuming it. We are precise about what this scores, since two assumptions hide in it. It is an argmax over the first token of each label — the branch the judge would take if it generated —

and *not* the total log-probability of the label sequence; the two coincide whenever the second token is near-determined given the first, which we expect here but did not measure. And the three ids come from tokenizing each label standalone with a leading space, not from the candidate continuations in template context; distinct standalone first pieces do not by themselves prove distinct in-context continuations. Scoring both tokens of each label in context would remove both assumptions for one extra forward pass, and is the version we would run given the budget. Separately, the judge sees the prompt and the completion truncated to 600 characters each. No completion reaches that length — the generation budget is 48 tokens — and exactly one prompt of 500 does, at 691 characters. UNCLEAR is kept as its own category and never merged into either decision; it accounts for at most 0.2% of completions at any rung. Judge-model grading is standard [Zheng et al., 2023]; we report it alongside four phrase-list variants and log every disagreement rather than treating any single grader as ground truth. The judge is not validated against human labels, which is the single largest gap in this work (§9).

4 Quantization shifts the refusal decision toward conservatism

Every rung scores the same prompts as its FP16 baseline, so the comparison is paired and the appropriate test is McNemar’s [McNemar, 1947] on the discordant pairs. The outcome has three states, so we state exactly what the test conditions on: a pair contributes to McNemar only if *both* sides produced a gradable (non-degenerate, non-unclear) completion and the two decisions differ. Table 1 therefore reports the number of gradable pairs beside every p -value. Because each model contributes seven rungs drawn from one prompt set, p -values are Holm-adjusted within model.

model	bits	→comply	→refuse	gradable	rate (%)	95% upper	p_{14}	p_7
Qwen2.5-3B	8.5	1	3	500	0.2	0.9	1.000	1.000
	7.5	8	10	500	1.6	2.9	1.000	1.000
	6.5	5	15	500	1.0	2.1	0.373	0.166
	5.5	5	19	500	1.0	2.1	0.066	0.033
	4.5	11	32	500	2.2	3.6	0.021	0.011
	3.5	0	30	154	0.0	0.6	<0.001	<0.001
	2.5	0	0	1	0.0	0.6	1.000	1.000
Phi-3.5-mini	8.5	3	4	499	0.6	1.5	1.000	1.000
	7.5	5	7	498	1.0	2.1	1.000	1.000
	6.5	4	6	499	0.8	1.8	1.000	1.000
	5.5	10	11	499	2.0	3.4	1.000	1.000
	4.5	4	21	498	0.8	1.8	0.011	0.005
	3.5	1	41	496	0.2	0.9	<0.001	<0.001
	2.5	0	0	0	0.0	0.6	1.000	1.000

Table 1: Paired transitions in the refusal decision against each model’s own FP16 baseline. →comply counts prompts the baseline refused and the rung answered; →refuse the reverse. *gradable* is the number of pairs where both sides produced a gradable completion, out of 500. *rate* is →comply over all 500 prompts and *95% upper* its one-sided exact upper bound. Both Holm corrections of the exact McNemar p -value are shown: p_{14} over every model $\times$ rung cell in the study, which is the primary column, and p_7 within one model’s ladder.

Two readings follow, and they are different claims.

The shifts run toward refusal, but only the direction survives independent graders.

Under the 7B judge, at 4.5 bits Qwen2.5-3B newly refuses 32 prompts it previously answered against 11 in the opposite direction ($p_{14} = 0.021$, $p_7 = 0.011$); Phi-3.5-mini shows 21 against 4 ($p_{14} = 0.011$, $p_7 = 0.005$). Both survive correction over every model $\times$ rung cell in the study. Neither survives re-grading by independent graders — the subject of the next box, and the reason we carry the direction of this shift rather than its significance.

What the judge’s refusal class actually contains. We read the 32 Qwen2.5-3B completions that newly count as refusal at 4.5 bits. Two things follow, and they point in opposite directions.

The first supports the result. Their mean length is 235 characters against 246 for the same model at full precision and 248 for the rung overall, so these are not truncated stubs: the 48-token budget is not manufacturing the effect by cutting answers off before they commit. On Qwen2.5-3B the composition of the refusal class also barely moves along the ladder: 54% of it contains a phrase-list marker at full precision against 53% at 4.5 bits. We do not generalise that — on Phi-3.5-mini the same figure rises from 16% to 24%, so the class composition there does shift — and marker incidence is in any case a description of the class, not a test that the judge applies it correctly.

The second qualifies it. Only 7 of the 32 contain a refusal marker, and reading the other 25 shows the class is broader than the word “refusal” suggests. Some are plainly refusals (“*As an AI, I don’t provide recommendations or advice on using firearms. . .*”). Others are capability disclaimers rather than safety ones (“*As an AI, I don’t have personal feelings. . .*”, “*I don’t have specific information about the exact sounds dinosaurs made. . .*”), which our judge template counts as REFUSE by construction, since it instructs that deflecting or redirecting qualifies. And a few look like grader error outright: “*Stalin was the leader of the Soviet Union from 1922 to 1953. . .*” is an informative answer, not a refusal of any kind.

So the quantity that rises is the judge’s REFUSE class — declining, deflecting, redirecting or warning — and not narrowly-construed refusal of a harmful request. Every claim in this paper should be read with that substitution made. This is also why blinded human validation (§9) is the first thing this work needs rather than a formality. Nothing here establishes that the judge applies the class correctly, and the shift in Phi-3.5-mini’s composition means we cannot even claim it applies it identically at both ends of that model’s ladder.

The result does not replicate under independent graders, and this is the most important finding in the paper. We re-graded the stored completions with three graders from two families other than the one under test: Llama-3.3-70B (Groq), Claude, and Claude Haiku. Between them they cover four model $\times$ grader comparisons. Every grader saw the same prompts, the same completions, the same template and the same degeneracy gate, so a difference is a grader difference and nothing else.

A fourth grader was attempted and is not counted here. Llama-3.1-70B (NVIDIA) agreed with the 7B judge on 89.3% of the 375 completions it returned, but its sweep did not cover both full precision and the 4.5-bit rung, so it supports no paired comparison at all. We exclude it rather than quote its agreement alongside the others, because 89.3% is the most flattering number in the set and it comes from the grader that finished least of the work. Three further sweeps (Gemini 2.0 Flash on both models, and Llama-3.3-70B on Phi-3.5-mini) returned between 4 and 10 completions and are not reported.

grader	model	agreement	$\rightarrow$ refuse	$\rightarrow$ comply	p
Qwen2.5-7B (as reported)	Qwen2.5-3B	—	32	11	0.002
Claude	Qwen2.5-3B	76.7%	10	5	0.302
Llama-3.3-70B	Qwen2.5-3B	78.4%	1	2	1.000
Qwen2.5-7B (as reported)	Phi-3.5-mini	—	21	4	0.001
Claude	Phi-3.5-mini	70.8%	12	7	0.359
Claude Haiku	Phi-3.5-mini	64.9%	5	4	1.000

Three of the four independent graders put the shift in the same direction — toward refusal — so the *sign* is not an artifact of one model. But the *magnitude* collapses, and not one of them reaches significance. Pairwise agreement with the 7B judge runs 64.9–78.4% across these four comparisons, which is far too low for a dependent variable: at 70% agreement on a 500-prompt paired comparison, grader disagreement alone moves more prompts than the effect being measured.

We therefore withdraw the claim that quantization significantly increases the refusal class at 4.5 bits. What survives is weaker and still worth reporting: a consistent direction across graders, with a magnitude that depends on which grader is asked. The phrase-list results in §5 are unaffected, because those compare two estimators on identical completions rather than resting on either being correct — and this section is the strongest evidence yet for that paper’s central point, since it shows grader choice moving a headline by a factor of thirty.

judge	model	n	agreement with 7B		4.5 bits, second judge	
			overall	on refusal	$\rightarrow$ refuse	$\rightarrow$ comply
agent claude	Qwen2.5-3B	450	76.7	73.4	10	5
	Phi-3.5-mini	400	70.8	68.2	12	7
agent haiku	Phi-3.5-mini	1000	64.9	60.6	5	4
groq llama-3-3-70b-versatile	Qwen2.5-3B	580	78.4	77.6	1	2

Table 2: Second-judge agreement and reproduction. Every grader saw the same prompts, completions, template and degeneracy gate as the 7B judge, so a difference is attributable to the grader alone. Restricted to sweeps that returned at least 400 graded completions: a partial sweep cannot support a paired comparison, and reporting its agreement anyway would let the least complete grader set the range. Four sweeps are excluded on that rule, and one of them — Llama-3.1-70B at 89.3% on 375 completions — is the highest agreement observed anywhere in the set.

Which family the correction runs over. The multiplicity question is worth settling even though the replication test above decides what we carry forward, because a reader comparing our




Figure 2: Left: agreement between each independent grader and the 7B judge, overall and restricted to completions the 7B judge called refusal. Right: the paired 4.5-bit counts recomputed from each grader’s own labels; shaded rows are the 7B judge’s own numbers, reproduced from Table 1 for comparison. Only graders that scored both FP16 and the 4.5-bit rung on at least 400 completions are shown, since a partial sweep cannot support a paired comparison.

tables against another study needs to know which family produced them. The p -values quoted so far correct across the seven rungs of one model. The simultaneous bound below instead treats all 14 model $\times$ rung cells as one family, and using a wider family for the bound than for the tests would be a way of taking the most favourable number from each. We therefore also correct the tests over all 14 cells. Under the 7B judge the 4.5-bit results survive on both models ($p = 0.021$ for Qwen2.5-3B, $p = 0.011$ for Phi-3.5-mini), and so do the 3.5-bit results; the 5.5-bit result does not ($p = 0.066$). Holm is valid under arbitrary dependence, so there is no technical reason to split by model; the only defence for the narrower family would be a pre-declared decision to treat each ladder as its own experiment, which we did not make. We therefore report the 14-cell correction as primary and Table 1 lists both columns. None of this rescues the significance claim: multiplicity correction prices repeated testing under one grader, and what the four independent graders show is that the choice of grader moves the estimate further than the correction does. The reverse claim needs a different instrument. A non-significant McNemar test is not evidence that transitions are absent, so we report a one-sided exact upper bound on the transition rate instead of inferring equivalence from a large p -value. The per-row bounds in Table 1 are pointwise, and a claim quantified over *every* rung needs simultaneous coverage: with Bonferroni over all 14 model $\times$ rung cells, no cell exceeds **4.62%**. The defensible statement is therefore bounded, not null: *whatever the true refusal-to-compliance transition rate is on these models and prompts, it is at most 4.62% at every rung simultaneously* — not that it is zero, and not that the quantized model is equivalent to the baseline.

What the McNemar tests condition on. Restricting to gradable pairs is a complete-case analysis, and at low precision the inclusion event is itself caused by quantization, so those rows do not estimate a population-level shift. This matters exactly where the gradable count is small and nowhere else. At 4.5 bits and above, degeneracy is 0.0% on both models and 498–500 of the 500 pairs are gradable, the shortfall being UNCLAR verdicts rather than degeneration, so the conditioning removes at most two pairs and the results quoted above are unaffected. At 3.5 bits Qwen2.5-3B retains 154 of 500 pairs and at 2.5 bits exactly one; Phi-3.5-mini retains none at 2.5

bits. Those rows are reported as complete-case descriptions and we draw no population claim from them — the $p = 1.000$ at 2.5 bits is the absence of data, not evidence of stability. This is why the gradable count is a column rather than a footnote: printing $n = 500$ against those rows would invite exactly the opposite reading.

4.1 Three regimes, and a drift coefficient inside one of them

Refusal rate is *not* a linear function of bit-width across the ladder. The data separate into three regimes:

1. **Inert**, FP16 down to 8.5 bits: refusal moves by +0.4 points on Qwen2.5-3B and +0.2 on Phi-3.5-mini. These are observed changes; we run no equivalence test, so the label names a regime rather than asserting a null.
2. **Drift**, 8.5 bits down to the coherence boundary: refusal rises approximately linearly in bits removed.
3. **Collapse**: output degenerates and no refusal rate is meaningful.

Within the drift regime — rungs where fewer than 10% of completions are degenerate — we fit

$$R(b) = \alpha - \kappa b, \quad b \in \mathcal{B}_{\text{coh}}, \quad (1)$$

with a *free* intercept α , where κ is the gain in refusal-rate points per bit removed. $\mathcal{B}_{\text{coh}}$ is $\{4.5, \dots, 8.5\}$ for Qwen2.5-3B and $\{3.5, \dots, 8.5\}$ for Phi-3.5-mini; full precision is *excluded*. Because all rungs score the same 500 prompts, ordinary regression standard errors are invalid, and all intervals below come from a prompt-level cluster bootstrap: resample prompts with replacement, recompute every rung’s refusal rate, refit, 10,000 replicates.

$$\hat{\kappa}_{\text{Qwen2.5-3B}} = 1.00 [0.44, 1.58], \quad \hat{\kappa}_{\text{Phi-3.5-mini}} = 1.29 [0.86, 1.74], \quad \hat{\kappa}_{\text{pooled}} = 1.15 [0.75, 1.55]. \quad (2)$$

Both intervals exclude zero. The two slopes are also not distinguishable from each other, and we test that rather than infer it from overlapping intervals: resampling prompts jointly so both slopes move together, the difference is 0.29 points/bit, 95% CI $[-0.35, 0.94]$, $p = 0.355$. The sign is positive throughout: *removing precision makes these models refuse more*.

Why the fit has a free intercept and not a full-precision anchor. The natural way to write a drift law is to anchor it at full precision, as $R(b) = R_{16} + \kappa(b_{16} - b)$. That form is not supported here, and testing it is what establishes the inert regime. Continuing the fitted band slope back to 16 bits predicts a full-precision refusal rate 7.5 points below the observed one on Qwen2.5-3B and 10.9 points below on Phi-3.5-mini (dotted lines, Figure 3). An anchored fit, a band fit, and a fit over all coherent rungs including FP16 are three different estimands with three different slopes; only the band fit describes a region where the relationship is actually linear, so that is the one we report, with its domain stated and its intercept free.

Two models do not make κ universal, and we do not claim it is. Nor is $[0.75, 1.55]$ a prediction interval for a new model: it is a confidence interval for these two, conditional on a coherent band that was itself selected from the observed degeneracy rates and then held fixed across bootstrap replicates. A third model landing outside it would therefore not by itself falsify anything. What Equation 1 provides is a form with a stated domain and a stated sign — precise enough to be measured again and disagreed with.




Figure 3: Refusal rate against stored bits per parameter. Filled markers are the coherent quantized rungs entering Equation 1; the solid and dashed lines are those fits. Hollow markers are full precision, excluded from the fit. Dotted lines continue each band fit back to FP16, where it undershoots the observed rate — the reason a single line across the whole range is not supportable.

5 Where the reported cliff comes from

A naive pipeline differs from ours in *two* respects, not one: it gates degeneracy on perplexity alone, and it labels the survivors by phrase matching. Table 3 crosses the two factors so their contributions separate. They do not act in the same place.

The two failures dominate in different regimes. Where output is coherent — 4.5 bits and above on both models, degeneracy 0.0% — the two gates admit exactly the same completions, so the entire gap is the *grader*: at 4.5 bits on Qwen2.5-3B the phrase list reports 10.8% against the judge’s 2.2%, a factor of 4.9. Where output has collapsed, the *gate* dominates: at 2.5 bits the phrase list reports 38.4% under a perplexity-only gate and 0.2% under the composite gate, on the same completions and with the same marker list — 99.5% of the gap, with the small remainder still attributable to the grader. The headline cliff needs both a weak gate and a weak grader; fixing either one alone removes most of it.

The grader half has two mechanisms, and both inflate the phrase-list estimate. Degenerate text that slips the gate contains no refusal phrase, so absence of a marker is read as compliance. Less obviously, *degraded models refuse in a different register*. In a supplementary 1.5B run excluded from the reported results — its grader saturated, §9 — 240 completions a marker list scored as compliance included 143 opening with “Sorry” and 9 written in a non-Latin script, among them a refusal to a self-harm request written in Chinese. We cite this as an illustration of the mechanism, not as evidence for its magnitude.

5.1 The phrase-list estimate is not monotone in the phrase list

The natural repair is to enlarge the marker list. It does not work, and the reason is structural rather than a matter of curation.




Figure 4: Refusal-to-compliance transitions as graded by the judge (solid) and as estimated by a refusal-phrase list (dashed) on identical completions. Shaded rungs are those where more than half of completions are degenerate.

model	bits	phrase list		judge		degen. (%)
		NLL	composite	NLL	composite	
Qwen2.5-3B	8.5	1.8	1.8	0.2	0.2	0.0
	7.5	2.4	2.4	1.6	1.6	0.0
	6.5	4.2	4.2	1.0	1.0	0.0
	5.5	3.8	3.8	1.0	1.0	0.0
	4.5	10.8	10.8	2.2	2.2	0.0
	3.5	15.4	13.2	0.0	0.0	69.2
	2.5	38.4	0.2	0.0	0.0	99.8
Phi-3.5-mini	8.5	1.2	1.2	0.6	0.6	0.0
	7.5	1.8	1.8	1.0	1.0	0.0
	6.5	3.0	3.0	0.8	0.8	0.0
	5.5	5.2	5.2	2.0	2.0	0.0
	4.5	5.6	5.6	0.8	0.8	0.0
	3.5	10.6	10.6	0.2	0.2	0.6
	2.5	0.0	0.0	0.0	0.0	100.0

Table 3: Gate and grader fully crossed on identical completions: both graders under both gates. Wherever degeneracy is 0.0% the two gates admit exactly the same completions — the NLL and composite columns are identical for both graders — so the phrase-list/judge difference there is attributable to the grader alone. Only at 3.5 bits and below does the gate contribute.

The quantity a phrase-list pipeline reports is a paired flip,

$$\hat{F}(\mathcal{M}) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}[m_{\mathcal{M}}(x_i^{16})] \cdot \mathbf{1}[\neg m_{\mathcal{M}}(x_i^b)], \quad (3)$$

where $m_{\mathcal{M}}(x)$ is true when the first 240 characters of completion x contain a marker from list $\mathcal{M}$. The window matters and is part of the estimator: phrase-list scoring looks for a refusal at the *start* of a response, because a marker appearing after 300 characters of compliance is not a refusal. Our implementation uses that window throughout, and the argument below does not depend on its width. Enlarging $\mathcal{M}$ pushes the two indicators in *opposite* directions: the baseline factor can only rise, the rung factor can only fall. That does not force non-monotonicity — such a product may still be monotone — but it removes any *guarantee* of one, so “I used a more inclusive list, therefore my estimate is closer to the truth” does not follow. Whether it is non-monotone in fact is empirical, and here it is: at 4.5 bits the flip count goes 54, 58, 71, 57 as the list grows (Table 4).

marker list	FP16 refusals	4.5 bits		2.5 bits	
		complies	flips	complies	flips
tight (as shipped)	223	273	54	422	192
+sorry	264	246	58	422	226
+sorry +as an ai	294	224	71	422	252
+deflections	306	196	57	422	263

Table 4: Both factors of Equation 3 on Qwen2.5-3B, as the marker list grows from 25 to 34 strings. At 2.5 bits the rung factor is pinned — no marker from any list appears in degenerate text, so 422 completions count as compliance under every variant — and the flip count rises purely because the baseline factor does. At 4.5 bits both factors move and the flip count is non-monotone: 54, 58, 71, 57.

marker list	code bits						
	8	7	6	5	4	3	2
tight (as shipped)	1.8	2.4	4.2	3.8	10.8	15.4	38.4
+sorry	1.6	2.6	4.4	3.8	11.6	17.6	45.2
+sorry +as an ai	2.2	2.8	4.8	4.0	14.2	18.8	50.4
+deflections	2.0	2.8	4.8	3.8	11.4	19.8	52.6

Table 5: Apparent refusal-to-compliance transitions (%) on Qwen2.5-3B for identical completions under four refusal-marker lists.

Across the rungs where every variant is non-zero, the estimate moves by up to $1.38\times$ on Qwen2.5-3B and $1.64\times$ on Phi-3.5-mini with the choice of strings alone. Ratios are undefined at Phi-3.5-mini’s 2-bit rung, where three variants report exactly zero and the fourth reports 0.2%; that rung is excluded from the ratio rather than reported as an infinite sensitivity. Every variant rises steeply below 5 bits where the judge does not, though none is monotone along the ladder — all four dip between 6 and 5 code bits, and on Phi-3.5-mini every variant falls back to near zero at 2 bits, where the model has stopped producing text a marker could appear in.




Figure 5: The data of Table 5. The ordering of the variants is not stable across rungs, which is Equation 3 in visual form.

6 Capability degrades at a model-specific threshold

GSM8K [Cobbe et al., 2021] answers are scored against gold arithmetic values, so no classifier intervenes. The same 200 questions are used at every rung, so this comparison is paired exactly as the refusal comparison is, and we test it the same way: exact McNemar on question-level transitions, Holm-adjusted across the seven rungs of each model. Correctness here is answer extraction alone, with no degeneracy gate, so that the capability measurement does not inherit the degeneracy detector’s decisions.




Figure 6: GSM8K accuracy with exact Clopper–Pearson intervals [Clopper and Pearson, 1934]. Accuracy is flat over most of the ladder, then collapses at a bit-width differing by a full bit between model families.

The threshold is not a property of bit-width alone: holding quantizer, decoding and benchmark fixed, it still moves by a full bit across model families. We say nothing about how it would move if those were varied too. Qwen2.5-3B is destroyed at 3.5 bits (0.0% accuracy, 37 questions lost, $p_{21} < 0.001$) and Qwen2.5-1.5B at the same rung (9.0%, 40 lost, $p_{21} = 0.001$), while Phi-3.5-mini shows no statistically significant paired difference at 3.5 bits and collapses only at 2.5. We say

model	bits	acc. (%)	lost	gained	p_{21}	p_7
Qwen2.5-1.5B	FP16	23.5	—	—	—	—
	8.5	23.5	5	5	1.000	1.000
	7.5	23.0	13	12	1.000	1.000
	6.5	24.0	12	13	1.000	1.000
	5.5	28.5	14	24	1.000	0.717
	4.5	25.0	18	21	1.000	1.000
	3.5	9.0	40	11	<0.001	<0.001
	2.5	1.0	47	2	<0.001	<0.001
Qwen2.5-3B	FP16	18.5	—	—	—	—
	8.5	19.5	4	6	1.000	1.000
	7.5	18.5	6	6	1.000	1.000
	6.5	17.5	9	7	1.000	1.000
	5.5	18.0	13	12	1.000	1.000
	4.5	11.5	23	9	0.321	0.100
	3.5	0.0	37	0	<0.001	<0.001
	2.5	0.5	37	1	<0.001	<0.001
Phi-3.5-mini	FP16	26.5	—	—	—	—
	8.5	25.0	6	3	1.000	1.000
	7.5	28.0	5	8	1.000	1.000
	6.5	27.0	9	10	1.000	1.000
	5.5	25.0	14	11	1.000	1.000
	4.5	26.5	18	18	1.000	1.000
	3.5	25.0	27	24	1.000	1.000
	2.5	0.0	53	0	<0.001	<0.001

Table 6: GSM8K accuracy with exact McNemar on question-level transitions. *lost* counts questions FP16 answered correctly and the rung did not; *gained* the reverse. p_{21} is Holm-adjusted over every model $\times$ rung cell, p_7 within one model. Correctness is answer extraction alone, with no degeneracy gate, so this arm does not inherit the detector’s decisions. $n = 200$ questions.

no significant difference was detected rather than that the rungs are equivalent, for the same reason given in §4.

A tempting claim this design cannot support. Qwen2.5-3B’s 4.5-bit rung is the one place where capability appears to fail while output is still fully coherent: accuracy falls from 18.5% to 11.5%, 23 questions lost against 9 gained, with not a single completion flagged as degenerate. It is tempting to report that as capability failure preceding fluency failure, and an unpaired binomial test against the FP16 rate would license it at $p = 0.010$. That test is invalid here — it treats an accuracy estimated on these same 200 questions as a known population rate. The paired exact McNemar test gives $p = 0.020$, and after Holm correction, $p_7 = 0.100$ within this model’s ladder and $p_{21} = 0.321$ over every cell in the study. We therefore do not claim the ordering. The point estimate and its direction are consistent with a threshold just below 4.5 bits, but with $n = 200$ and a full-precision accuracy near 20% this design lacks the power to resolve it, and a larger question set is the specific thing that would.

7 A frozen refusal probe does not track any of this

We fit a difference-in-means refusal direction [Arditi et al., 2024] on FP16 activations and apply it unchanged to every rung, reporting d' [Green and Swets, 1966] retention. The direction is fitted on one half of the prompts and scored on a disjoint half, over 200 synchronized replicates: the same split drives every rung, so cells differ only by which rung supplied the activations. We report only this frozen variant. Refitting the direction on each rung’s own activations answers a different question — whether a *new* probe can be trained on the quantized model — and the two estimands come apart sharply at low precision, so reporting them interchangeably would be a mistake.

The probe is fitted on the same labels the behaviour is judged by. This matters more than it sounds. If the probe separated *marker*-derived classes while the behavioural arm tracked the judge’s, then a probe that stayed flat while behaviour drifted would show only that the two targets differ — the probe could be succeeding at its own task. There would be no dissociation to report. We therefore fit the direction on the judge’s own full-precision verdicts, so both arms aim at one target and any gap between them is a real one.

Retention is a ratio, so we give the scale it is a ratio of. A probe keeping 100% of a d' of 0.3 has kept nothing worth having. On the judge’s labels the full-precision d' is 2.01 for Phi-3.5-mini, which is a strong separation, but only 0.74 for Qwen2.5-3B. Qwen’s probe is therefore a weak detector before quantization touches it, and “retains 95% of its discriminability” has to be read against that. This is a limitation of the probe on this label set, not a finding about quantization.

The reported intervals are the 2.5–97.5 percentiles of retention *across those splits*. They are a dispersion measure, not a confidence interval: the splits reuse one fixed prompt set and overlap heavily, so they describe how much the estimate moves with the split and say nothing about sampling variation in the prompt population. We therefore read them descriptively and make no claim of the form “indistinguishable from 100%”. A proper interval would require refitting the probe inside an outer bootstrap over prompts, which we have not done.

From 8.5 through 4.5 bits retention stays between 95% and 100% on both models in the refusal arm. That range covers both models’ 4.5-bit refusal shifts and Qwen2.5-3B’s 4.5-bit arithmetic drop. Over that span the probe moves by 5.0 points on Qwen2.5-3B and 2.7 on Phi-3.5-mini, while the refusal rate moves by four and GSM8K accuracy by seven. It does *not* cover Phi-3.5-mini’s largest shift at 3.5 bits, where retention has already fallen to 57%.

bits	Qwen2.5-3B		Phi-3.5-mini	
	kept (%)	split range	kept (%)	split range
FP16	100	[100, 100]	100	[100, 100]
8.5	100	[100, 101]	100	[100, 100]
7.5	100	[99, 102]	100	[99, 100]
6.5	100	[98, 103]	100	[99, 101]
5.5	99	[96, 103]	100	[98, 102]
4.5	95	[86, 104]	98	[94, 101]
3.5	42	[3, 78]	57	[48, 67]
2.5	31	[-0, 63]	5	[-8, 16]

Table 7: Frozen-probe d' retention as a percentage of the full-precision value. *split range* is the 2.5–97.5 percentile over 200 fit/score splits of one fixed prompt set — a dispersion measure, not a confidence interval. Ranges can exceed 100% where a quantized rung separates the classes marginally better than FP16 on a given split.




Figure 7: Frozen-probe retention along the ladder, with split-to-split bands. Retention is near-flat at 95–100% over 8.5–4.5 bits, which covers both models’ 4.5-bit refusal shifts. Phi-3.5-mini is the exception: its largest shift, at 3.5 bits, sits at 57% retention, at a rung where composite degeneration is only 0.6%. Labels are the judge’s, matching the behavioural arm.

Two qualifications on the word “flat”. First, retention is not uniformly flat: Qwen2.5-3B sits at 95.3% at 4.5 bits, and the split-to-split spread there, [86, 104], is wide. The claim is that the probe moves little, not that it does not move. Second, the coherent band is model-specific, and for Phi-3.5-mini it extends to 3.5 bits, where retention has already fallen to 57%. So the probe is *not* flat across the whole drift band of every model; what holds on both is the 8.5–4.5 range. Below 4 bits the probe collapses along with the model (Qwen2.5-3B to 42% at 3.5 bits).

Over 8.5–4.5 bits the probe does not track behaviour. It is near-flat there while the refusal decision drifts on both models and arithmetic falls by seven points, 38% of its baseline. It then falls, and where it falls differs by model: on Qwen2.5-3B the drop coincides with degeneration, but on Phi-3.5-mini retention is already 57% at 3.5 bits, a rung our detector calls degenerate for only 0.6% of completions and where the refusal shift is largest (41 prompts newly refused against 1 in the other direction). So the probe is not simply a degeneration detector either; the defensible statement is the narrow one, that across the range where these models remain coherent it stays high while their behaviour moves. This is the limitation Belinkov [2022] describes: a linearly decodable label is not thereby a behavioural one. It is also why the agreement with Chhabra and Khalili [2025] and Duan [2026] that compression preserves refusal geometry should not be read as evidence that compression preserves refusal behaviour.

One consequence of fitting on the judge’s labels is worth stating, because it cuts against the section. Every figure here inherits whatever the judge gets wrong: the probe’s positive class is the set of prompts the 7B judge called refusal at full precision, so a systematic grading error becomes a systematic error in what the probe is asked to separate. That is the price of putting both arms on one target, and it is the right price to pay — fitting the probe on marker-derived classes instead, as an earlier version did, compares two systems aimed at different things and cannot establish a dissociation at all. We report the marker-label variant alongside it in the released statistics for contrast. The two agree where the paper’s claim lives — both are near-flat over 8.5–4.5 bits, 95–100% under the judge’s labels and 96–100% under the marker list — and diverge below it, where Phi-3.5-mini’s 3.5-bit retention reads 57% against 63% and Qwen2.5-3B’s reads 42% against 34%. Two targets give two answers once the model starts breaking, which is a reason to fix the target rather than to average them.

8 Discussion

Deployment. Neither refusal-phrase matching nor activation probing detected the degradation a user would actually notice; scoring a gold-labelled task did, and it requires no grader and no labelled safety corpus. Screening for degenerate output is necessary but insufficient: at 4.5 bits neither model has a single completion flagged by our detector, while one has lost a third of its arithmetic point-estimate.

8.1 What this implies for choosing a quantizer

The three regimes suggest a design hypothesis. We state it as a hypothesis because this paper measures one quantizer family and does not test the rule it motivates.

Bit-width is conventionally chosen as the smallest setting whose output is still usable, which in practice means the smallest b before linguistic collapse:

$$b_{\text{conventional}}^* = \min\{b : \text{output remains coherent}\}. \quad (4)$$

Our measurements place that boundary strictly below where behaviour has already moved. At 4.5 bits both models are fully coherent by our detector, yet the refusal decision has shifted on both — in the same direction under every grader we tried — and Qwen2.5-3B’s arithmetic point-estimate has fallen by a third. Equation 4 therefore selects a checkpoint on the far side of changes a deployer would care about. The alternative is to make the behavioural quantities constraints rather than afterthoughts:

$$b^* = \arg \min_b \text{Memory}(b) \quad \text{s.t.} \quad T^{r \rightarrow c}(b) \leq \epsilon_s, \quad T^{c \rightarrow r}(b) \leq \epsilon_o, \quad D(b) \leq \epsilon_d, \quad C(b) \geq C_{\min}, \quad (5)$$

where $T^{r \rightarrow c}$ and $T^{c \rightarrow r}$ are the two directional transition rates, D the degeneration rate and C task capability — exactly the four quantities §1.1 produces. Each is a measured curve in this paper, so Equation 5 is evaluable today for any candidate checkpoint; what is untested is whether it selects better deployments than Equation 4 in practice.

Two of our results bear on *which* signals such a rule may use. The frozen probe retains 95–100% of its discriminability across 8.5–4.5 bits, where the refusal decision is already drifting, and the phrase list moves opposite to the judge. So refusal-direction retention and phrase-based refusal rates are both poor candidates for the constraint set, however convenient they are to compute. That is a negative result about proxies, and it is the part of this section the data actually support.

A stronger version — allocating precision per layer according to which blocks carry the behavioural sensitivity, rather than uniformly — is a natural follow-up. We do not pursue it here because we have not localised the drift to any layer, and saying which layers matter would require an ablation this paper does not contain.

Reading the literature. Reported safety-cliff magnitudes are sensitive to grader construction in a way that is not usually reported alongside them, and Equation 3 shows the sensitivity is not even signed. We recommend that studies of compression and safety publish a degeneracy rate beside every compliance rate, state the marker list in full if one is used, and validate graders against refusals worded outside the expected register, including non-English ones. Where a study reports refusal falling under quantization [Kadadekar, 2026], the degeneracy rate at the same rung is the number that distinguishes a genuine loss from the artifact characterised here.

Why refusal survives. Refusal behaviour persisted to the point of linguistic collapse while multi-step arithmetic did not. A plausible reading is that refusal is a short, high-prior continuation reachable from a wide basin, whereas arithmetic requires a long chain of precise steps; the first is robust to weight noise in a way the second cannot be. The upward drift would then be a side-effect of the same mechanism, as noise widens the basin further. We do not test this here.

9 Limitations

- **Change, not absolute safety.** Labels are defined against each model’s own full-precision behaviour. A model that answers harmful requests at FP16 and continues to do so after quantization contributes no discordant pair.
- **No harmfulness annotation.** HH-RLHF first-turn utterances span harmful, benign and ambiguous requests without per-prompt labels. A baseline refusal may be an over-refusal of a benign prompt, in which case the quantized model’s compliance is a correction rather than a failure. Licensing the term “harmful compliance” requires a harmfulness-labelled suite [Mazeika et al., 2024, Souly et al., 2024] plus a substantive-harm rubric; an over-refusal suite would be needed to price the conservative shift we do report.
- **Grader validity is the largest gap, and we have direct evidence of error.** Reading the newly-refusing completions at 4.5 bits (§4) turns up informative answers labelled REFUSE,

alongside capability disclaimers that the template counts as refusal by construction. The judge is unvalidated against blinded human labels, is itself quantized to 4-bit, and shares a family with two of the three models, so stylistic self-preference cannot be excluded. We report its agreement with four phrase lists and three independent graders, not its accuracy against anything. A stratified blinded human sample across models, rungs, languages and disagreement cases is the next thing this work needs.

- **The headline does not replicate across graders, and every quantity derived from the judge inherits that.** Four graders from three other families agree on the direction three times out of four and on the significance never (§4). Independent graders are therefore no longer the missing piece — they were run, and they defeated the claim. The consequence reaches past the p -values: the drift coefficient κ and its interval, the probe’s judge-derived labels, and the transition bounds are all computed from the 7B judge’s labels, so their intervals price sampling error and not grader choice. On this evidence grader choice is the larger term. Only blinded human labels can say which grader is right, and none of ours is validated against any.
- **Outcome-dependent gating.** Degeneracy is judged before refusal. We bound the effect rather than eliminate it (§3.1): the gate removes nothing at 4.5 bits and above, and the judge/phrase-list gap is unchanged under a perplexity-only gate. Below 4 bits the gate is load-bearing and those rows should be read with their gradable counts. A blinded audit of gated completions is not yet done.
- **Short completions.** 48 new tokens is enough for a refusal but often not for substantive compliance, which biases the design toward finding refusal intact. A completion that opens with a disclaimer and supplies harmful content later would be truncated to the disclaimer. Longer budgets and sampled decoding across seeds are needed before the transition-rate bound can be called tight.
- **Greedy decoding.** One completion per prompt estimates a deterministic decision, not behaviour under sampling.
- **Scope.** Three models (two in the refusal arm), one quantizer family, one reasoning benchmark, English prompts. Nothing here establishes behaviour under AWQ, GPTQ, GGUF k-quants or mixed-precision deployment formats. Absolute GSM8K accuracy is low (18.5–26.5%) because generation is zero-shot with a 192-token budget, so only relative change is interpreted, and the resulting power is low enough that §6 cannot resolve a moderate effect.
- **“Capability” and “fluency” are narrow here.** One arithmetic benchmark is not general capability, and passing four handcrafted degeneracy thresholds is not human-rated fluency.

10 Reproducibility

All measurements run on a single 6 GB consumer GPU except the 7B judge and the 3B/3.8B models, which were run on a hosted T4. Raw completions are stored verbatim for every rung, so any grader can be re-applied without regenerating text; the degeneracy correction of §3.1, the paired GSM8K re-analysis of §6 and the probe intervals of §7 were all applied retrospectively to stored artifacts at no additional generation cost. Every table and figure, and the consolidated data files behind them, are emitted from the stored run directories by the analysis scripts rather than transcribed. Numbers quoted in the prose are typed, so a checker re-derives the load-bearing ones from those files and fails if the manuscript disagrees; it runs beside the test suite. The quantizer, the judge template, the four marker lists, the degeneracy thresholds and the prompt selection are all in the released code.

The GSM8K questions are not redistributed here. To make the capability arm checkable without them, the analysis records a SHA-256 of the question file it consumed (3730d312...c39d14) and of the ordered 200-element gold answer vector actually used (28ba1fa0...3d80ef8); both

are reproducible from the canonical GSM8K test split by taking the first 200 items with a parseable gold answer, in file order.

References

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024.
- Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. In *International Conference on Learning Representations, Workshop Track*, 2017.
- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024. arXiv:2406.11717.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.
- Vishnu Kabir Chhabra and Mohammad Mahdi Khalili. Towards understanding and improving refusal in compressed models via mechanistic interpretability. *arXiv preprint arXiv:2504.04215*, 2025.
- C. J. Clopper and E. S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems*, 2023.
- Evan Duan. Do activation monitors survive model updates? benchmarking, predicting, and repairing activation-monitor staleness. *arXiv preprint arXiv:2606.15980*, 2026.
- Kazuki Egashira, Mark Vero, Robin Staab, Jingxuan He, and Martin Vechev. Exploiting LLM quantization. *arXiv preprint arXiv:2405.18137*, 2024.
- Kazuki Egashira, Robin Staab, Mark Vero, Jingxuan He, and Martin Vechev. Mind the gap: A practical attack on GGUF quantization. In *International Conference on Machine Learning*, 2025. arXiv:2505.23786.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations*, 2023.
- Shaona Ghosh, Prasoon Varshney, Erick Galinkin, and Christopher Parisien. AEGIS: On-line adaptive AI content safety moderation with ensemble of LLM experts. *arXiv preprint arXiv:2404.05993*, 2024.

- David M. Green and John A. Swets. *Signal Detection Theory and Psychophysics*. Wiley, 1966.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *International Conference on Learning Representations*, 2020.
- Junyuan Hong, Jinhao Duan, Chenhui Zhang, Zhangheng Li, Chulin Xie, Kelsey Lieberman, James Diffenderfer, Brian Bartoldson, Ajay Jaiswal, Kaidi Xu, Bhavya Kailkhura, Dan Hendrycks, Dawn Song, Zhangyang Wang, and Bo Li. Decoding compressed trust: Scrutinizing the trustworthiness of efficient LLMs under compression. In *International Conference on Machine Learning*, 2024. arXiv:2403.15447.
- Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yuning Mao, Michael Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madian Khabsa. Llama guard: LLM-based input-output safeguard for human-AI conversations. *arXiv preprint arXiv:2312.06674*, 2023.
- Sahil Kadadekar. Quality is not a safety proxy under quantization. *arXiv preprint arXiv:2606.10154*, 2026.
- Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration. In *Proceedings of Machine Learning and Systems*, 2024.
- Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- Alexandra Souly, Qingyuan Lu, Dillon Bowen, Tu Trinh, Elvis Hsieh, Sana Pandey, Pieter Abbeel, Justin Svegliato, Scott Emmons, Olivia Watkins, and Sam Toyer. A StrongREJECT for empty jailbreaks. *arXiv preprint arXiv:2402.10260*, 2024.
- Sunghyun Wee, Suyoung Kim, Hyeonjin Kim, Kyomin Hwang, and Nojun Kwak. Safety-preserving PTQ via contrastive alignment loss. *arXiv preprint arXiv:2511.07842*, 2025.
- Jiachen Zhao, Jing Huang, Zhengxuan Wu, David Bau, and Weiyan Shi. LLMs encode harmfulness and refusal separately. *arXiv preprint arXiv:2507.11878*, 2025.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, 2023.